On the usage of the `geepack`

Søren Højsgaard and Ulrich Halekoh

geepack version 1.3.11 as of 2024-06-06

Contents

1 Introduction

This note contains a few extra examples. We illustrate the usage of a the `waves` argument and the `zcor` argument together with a fixed working correlation matrix for the `geeglm()` function.

1.1 Citing `geepack`

The primary reference for the `geepack` package is

> Halekoh, U., Højsgaard, S., Yan, J. (2006) *The R Package geepack for Generalized Estimating Equations (2006)* Journal of Statistical Software `https://www.jstatsoft.org/article/view/v015i02`

```
> library(geepack)
> citation("geepack")
```

To cite geepack in publications use:

```
  Højsgaard, S., Halekoh, U. & Yan J. (2006) The R Package geepack for
  Generalized Estimating Equations Journal of Statistical Software, 15,
  2, pp1--11

  Yan, J. & Fine, J.P. (2004) Estimating Equations for Association
  Structures Statistics in Medicine, 23, pp859--880.
```

Yan, J (2002) geepack: Yet Another Package for Generalized Estimating Equations R-News, 2/3, pp12-14.

To see these entries in BibTeX format, use 'print(<citation>, bibtex=TRUE)', 'toBibtex(.)', or set 'options(citation.bibtex.max=999)'.

If you use geepack in your own work, please do cite the above reference.

1.2   When do GEE's work best?

1. GEEs work best when you have relatively many relatively small clusters of about equal size in your data.

2. If all your clusters are of size one you should not use GEEs; if all clusters are of size one a GEE corresponds to a generalized linear model.

3. If you only have very few clusters (and in the extreme case only one cluster) you are likely to encounter numerical difficulties.

NOTICE: Care must be taken with respect to the order in which the clusters appear in the dataset. See Section 3 for details.

2   Simulating a dataset

To illustrate the usage of the `waves` argument and the `zcor` argument together with a fixed working correlation matrix for the `geeglm()` we simulate data suitable for a regression model.

```
> library(geepack)
> n_cluster <- 6
> n_time    <- 5
> set.seed(1213)
> timeorder <- rep(1:n_time, n_cluster)
> tvar      <- timeorder + rnorm(length(timeorder))
> idvar  <- rep(1:n_cluster, each=n_time)
> uuu    <- rep(rnorm(n_cluster), each=n_time) # A 'random intercept'
> yvar   <- 1 + 2 * tvar + uuu + rnorm(length(tvar))
> simdat <- data.frame(idvar, timeorder, tvar, yvar)
> head(simdat, 12)
```

```
   idvar timeorder       tvar       yvar
1      1         1   1 -0.1199294  0.7082247
2      1         2   3.0411260  7.1458370
3      1         3   0.5028488  3.3108124
4      1         4   1.7604888  3.8893380
5      1         5   6.1349016 14.2613890
6      2         1   1.2554963  3.4235073
7      2         2   0.6098986  3.9045993
8      2         3   2.8755870  6.0865324
9      2         4   3.9935618  9.7830956
10     2         5   6.0638870 14.0556776
```

```
11     3          1  1.8977871  5.6986661
12     3          2  3.2235643  6.9993705
```

Notice that clusters of data appear together in `simdat` and that observations are ordered (according to `timeorder`) within clusters.

We can fit a model with an AR(1) error structure as

```
> mod1 <- geeglm(yvar~tvar, id=idvar, data=simdat, corstr="ar1")
> mod1

Call:
geeglm(formula = yvar ~ tvar, data = simdat, id = idvar, corstr = "ar1")

Coefficients:
(Intercept)        tvar
  0.6314834   1.9908676

Degrees of Freedom: 30 Total (i.e. Null);  28 Residual

Scale Link:                     identity
Estimated Scale Parameters:   [1] 1.390467

Correlation:  Structure = ar1     Link = identity
Estimated Correlation Parameters:
    alpha
0.6521569

Number of clusters:   6   Maximum cluster size: 5
```

This works because observations are ordered according to time within each subject in the dataset.

3   Using the `waves` argument

If observatios were not ordered according to cluster and time within cluster we would get the wrong result:

```
> set.seed(123)
> simdatPerm <- simdat[sample(nrow(simdat)),]
> simdatPerm <- simdatPerm[order(simdatPerm$idvar),]
> head(simdatPerm)

   idvar timeorder        tvar        yvar
3      1         3   0.5028488   3.3108124
5      1         5   6.1349016  14.2613890
4      1         4   1.7604888   3.8893380
1      1         1  -0.1199294   0.7082247
2      1         2   3.0411260   7.1458370
10     2         5   6.0638870  14.0556776
```

Notice that in `simdatPerm` data is ordered according to subject but the time ordering within subject is random.

Fitting the model as before gives

```
> mod2 <- geeglm(yvar~tvar, id=idvar, data=simdatPerm, corstr="ar1")
> mod2

Call:
geeglm(formula = yvar ~ tvar, data = simdatPerm, id = idvar,
    corstr = "ar1")

Coefficients:
(Intercept)        tvar
  0.6802353   2.0048127


Degrees of Freedom: 30 Total (i.e. Null);  28 Residual

Scale Link:                       identity
Estimated Scale Parameters:   [1] 1.38735

Correlation:  Structure = ar1     Link = identity
Estimated Correlation Parameters:
    alpha
0.6294383


Number of clusters:    6   Maximum cluster size: 5
```

Likewise if clusters do not appear contigously in data we also get the wrong result (the clusters are not recognized):

```
> simdatPerm2 <- simdat[order(simdat$timeorder),]
> head(simdatPerm2)

   idvar timeorder         tvar        yvar
1      1         1 -0.1199294 0.7082247
6      2         1  1.2554963 3.4235073
11     3         1  1.8977871 5.6986661
16     4         1  2.1254099 5.4517051
21     5         1  0.5043048 0.7625234
26     6         1  1.3327271 2.0536935


> geeglm(yvar~tvar, id=idvar, data=simdatPerm2, corstr="ar1")

Call:
geeglm(formula = yvar ~ tvar, data = simdatPerm2, id = idvar,
    corstr = "ar1")

Coefficients:
(Intercept)        tvar
  0.6935408   1.9920382


Degrees of Freedom: 30 Total (i.e. Null);  28 Residual

Scale Link:                       identity
Estimated Scale Parameters:   [1] 1.386165


Correlation:  Structure = ar1     Link = identity
```

```
Estimated Correlation Parameters:
alpha
     0

Number of clusters:   30   Maximum cluster size: 1
```

To obtain the right result we must give the `waves` argument:

```
> wav <- simdatPerm$timeorder
> wav

 [1] 3 5 4 1 2 5 4 3 2 1 5 4 1 3 2 4 3 5 2 1 2 4 5 3 1 3 2 1 5 4

> mod3 <- geeglm(yvar~tvar, id=idvar, data=simdatPerm, corstr="ar1", waves=wav)
> mod3


Call:
geeglm(formula = yvar ~ tvar, data = simdatPerm, id = idvar,
    waves = wav, corstr = "ar1")

Coefficients:
(Intercept)        tvar
  0.6314834   1.9908676

Degrees of Freedom: 30 Total (i.e. Null);  28 Residual

Scale Link:                    identity
Estimated Scale Parameters:   [1] 1.390467

Correlation:  Structure = ar1    Link = identity
Estimated Correlation Parameters:
    alpha
0.6521569

Number of clusters:    6   Maximum cluster size: 5
```

4   Using a fixed correlation matrix and the `zcor` argument

Suppose we want to use a fixed working correlation matrix:

```
> cor.fixed <- matrix(c(1    , 0.5  , 0.25, 0.125, 0.125,
+                        0.5  , 1    , 0.25, 0.125, 0.125,
+                        0.25 , 0.25 , 1   , 0.5  , 0.125,
+                        0.125, 0.125, 0.5 , 1    , 0.125,
+                        0.125, 0.125, 0.125, 0.125, 1    ), 5, 5)
> cor.fixed


      [,1]  [,2]  [,3]  [,4]  [,5]
[1,] 1.000 0.500 0.250 0.125 0.125
[2,] 0.500 1.000 0.250 0.125 0.125
```

```
[3,] 0.250 0.250 1.000 0.500 0.125
[4,] 0.125 0.125 0.500 1.000 0.125
[5,] 0.125 0.125 0.125 0.125 1.000
```

Such a working correlation matrix has to be passed to geeglm() as a vector in the zcor argument. This vector can be created using the fixed2Zcor() function:

```
> zcor <- fixed2Zcor(cor.fixed, id=simdatPerm$idvar, waves=simdatPerm$timeorder)
> zcor

 [1] 0.125 0.500 0.250 0.250 0.125 0.125 0.125 0.125 0.125 0.500 0.125 0.125
[13] 0.125 0.125 0.500 0.125 0.125 0.250 0.250 0.500 0.125 0.125 0.125 0.125
[25] 0.125 0.500 0.125 0.250 0.500 0.250 0.500 0.125 0.125 0.125 0.125 0.250
[37] 0.250 0.125 0.125 0.500 0.125 0.125 0.250 0.500 0.125 0.500 0.125 0.125
[49] 0.125 0.250 0.250 0.250 0.125 0.500 0.500 0.125 0.125 0.125 0.125 0.125
```

Notice that zcor contains correlations between measurements within the same cluster. Hence if a cluster contains only one observation, then there will be generated no entry in zcor for that cluster. Now we can fit the model with:

```
> mod4 <- geeglm(yvar~tvar, id=idvar, data=simdatPerm, corstr="fixed", zcor=zcor)
> mod4

Call:
geeglm(formula = yvar ~ tvar, data = simdatPerm, id = idvar,
    zcor = zcor, corstr = "fixed")

Coefficients:
(Intercept)        tvar
  0.6716075   1.9854861

Degrees of Freedom: 30 Total (i.e. Null);  28 Residual

Scale Link:                     identity
Estimated Scale Parameters:  [1] 1.388038

Correlation:  Structure = fixed    Link = identity
Estimated Correlation Parameters:
alpha:1
      1

Number of clusters:   6   Maximum cluster size: 5
```